




UNIVERSITAT
ROVIRA I VIRGILI



AWS Academy

Distributed Systems

Introduction to AWS S3

► What is AWS S3?

- Amazon Simple Storage Service (S3) is a highly scalable, durable, and secure object storage service.
- It is used to store and retrieve any amount of data at any time.

► Key Features:

- **Scalability:** Automatically scales storage capacity.
- **Durability & Availability:** Designed for 99.999999999% (11 9's) durability.
- **Security:** Supports encryption, access management, and compliance.
- **Cost-Effective:** Pay only for the storage you use.

► Use Cases:

- Data backup and archiving.
- Hosting static websites.
- Storing large amounts of unstructured data.

Introduction to AWS Lambda

► What is AWS Lambda?

- AWS Lambda is a serverless compute service that allows you to run code without provisioning or managing servers.
- It executes code in response to events, such as S3 uploads, API Gateway calls, or scheduled events.

► Key Benefits:

- **No Server Management:** Automatically scales with demand.
- **Event-Driven:** Code is triggered by specific events.
- **Cost Efficiency:** You pay only for the compute time you consume.
- **Flexible:** Supports multiple programming languages (e.g., Python, Node.js, Java).

► Common Use Cases:

- Real-time file processing (e.g., processing CSV uploads in S3).
- Backend for mobile or web applications.
- Automation of operational tasks.

Introduction to AWS EC2

▶ What is AWS EC2?

- ▶ Amazon Elastic Compute Cloud (EC2) provides scalable virtual servers in the cloud.
- ▶ It gives you full control of your computing resources and allows you to run applications on a virtual machine.

▶ Key Features:

- ▶ **Flexible Instance Types:** Choose from a variety of instance types optimized for different workloads.
- ▶ **Scalability:** Scale up or down based on demand.
- ▶ **Control:** Full access to the operating system, storage, and network.
- ▶ **Integration:** Easily integrates with other AWS services like S3, RDS, and more.

▶ Common Use Cases:

- ▶ Hosting web applications and microservices.
- ▶ Running batch processing jobs.
- ▶ Deploying legacy applications that require complete OS control.

Conceptual Overview

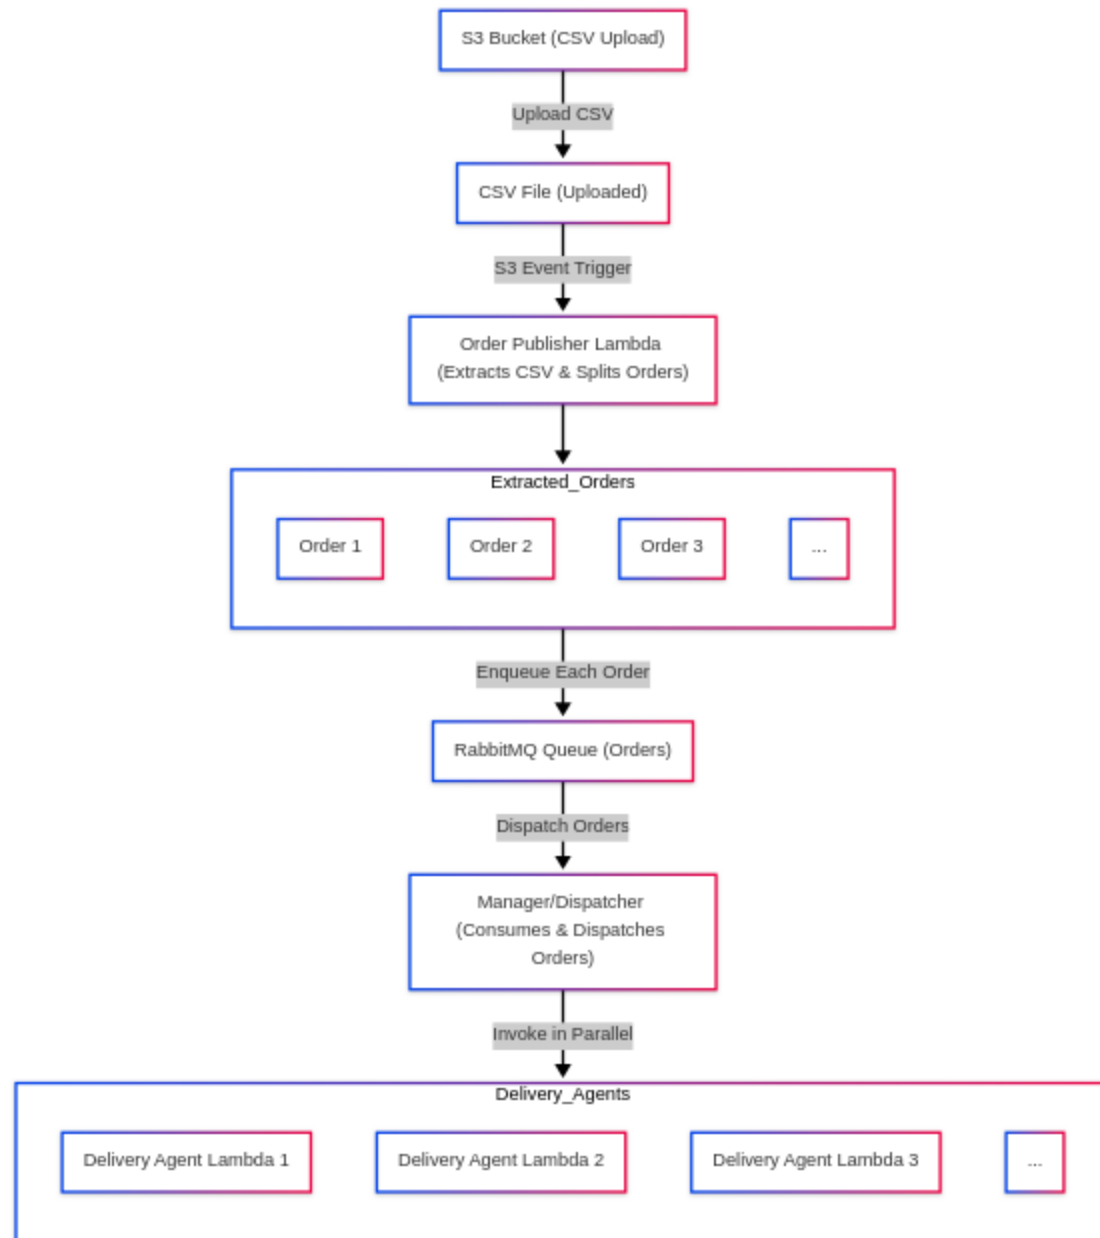
Distributed Order Processing System

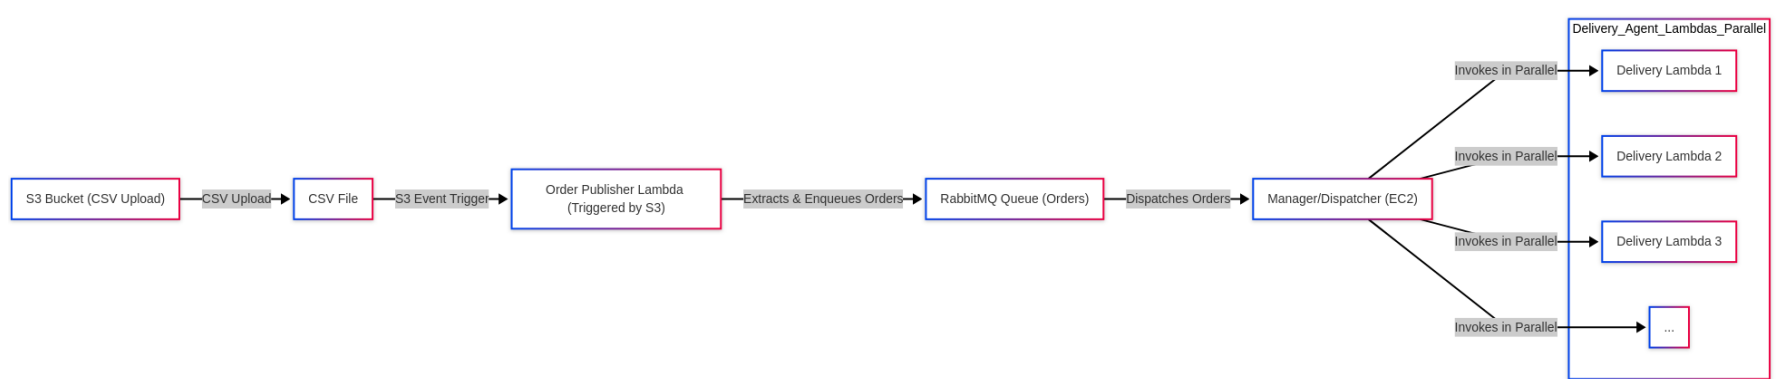
► Objective

- Simulate the operation of an order processing platform that receives customer orders, processes them, and assigns each order to a delivery agent. This system mimics real-world scenarios where orders must be handled concurrently by a fleet of couriers.

System Overview

- ▶ Build a distributed order processing system.
- ▶ The system starts by **receiving orders** via a CSV file uploaded to an **S3 bucket**.
- ▶ The uploaded CSV is processed by a Lambda function, which extracts each order and enqueues them in a **RabbitMQ queue**.
- ▶ A **Manager component** (or dispatcher) then takes the orders from the queue and **assigns each order to a different delivery agent**.
- ▶ Each delivery agent (simulated by a separate **Lambda function**) processes its order in parallel, mimicking concurrent delivery.





Key Steps

- ▶ **CSV Upload:** Orders are submitted in a CSV file.
- ▶ **Order Enqueueing:** A Lambda function reads the CSV, parses orders, and publishes them to a RabbitMQ queue.
- ▶ **Order Distribution:** A Manager component (running on an EC2 instance) subscribes to the queue and invokes a Delivery Simulator Lambda for each order.
- ▶ **Parallel Delivery:** Each invoked Lambda simulates order delivery concurrently.

Lab and Environment Setup

- ▶ Register on AWS Academy and log in to the LMS.
- ▶ Navigate to **Courses** → **All Courses** and open **AWS Academy Learner Lab**.
- ▶ In the module, click **Launch AWS Academy Learner Lab**.
- ▶ Start the lab with **Start Lab** and click the AWS icon once it turns green.

Creating an S3 Bucket

► Task:

- Create a new S3 bucket in your AWS environment.

► Purpose:

- This bucket will trigger the Order Publisher Lambda function when you upload the CSV file.

► Resource:

- [Creating an S3 Bucket Guide](#)

Order Publisher Lambda Function

- ▶ **Objective:**
 - ▶ Write a Python Lambda function (named **lambda_order_publisher**) that triggers on S3 events.
 - ▶ This function reads the uploaded CSV, extracts orders, and publishes them to a RabbitMQ queue.
- ▶ **Key Code Components:**
 - ▶ **Extract Event Data:** Read the bucket name and CSV file key from the S3 event.
 - ▶ **Read and Parse CSV:** Use boto3 to get the file, then `csv.DictReader` to parse it.
 - ▶ **Publish Orders:** Use the pika library to connect to RabbitMQ and send each order as a JSON message.
- ▶ **Tip:**
 - ▶ First test the Lambda without RabbitMQ code—just log the orders. Then add the RabbitMQ connection.

Order Publisher Lambda Function

- ▶ Important – Add Required Python Packages Using a Lambda Layer:
 - ▶ Go to the **Lambda Dashboard** in AWS.
 - ▶ Create a **new Lambda Layer**.
 - ▶ **Upload the file:** `layer_contents.zip`
(This ZIP file is already prepared and includes the necessary dependencies: pika and boto3, located inside a python/ folder.)
 - ▶ After creating the layer, **copy its ARN**.
 - ▶ Go to your `lambda_order_publisher` function.
 - ▶ In the "Layers" section, click **"Add a layer"**.
 - ▶ Select **"Provide a layer version ARN"** and **paste the ARN** you copied.
 - ▶ **Note:** For detailed steps on how to create a Lambda layer, **refer to the resources on the next slide**.

Order Publisher Lambda Function

► Important Notes:

- Make sure to configure your Lambda function's timeout to a sufficient value (for example, 30 seconds) in the AWS Lambda configuration. This ensures that the function has enough time to retrieve the file from S3, parse the CSV, establish a connection to RabbitMQ, and publish all the orders. A too-short timeout might cause your function to terminate before completing these tasks, leading to errors or incomplete processing.
- Remember to redeploy your Lambda function after each modification.

► Resources:

- [Packaging your layer content](#)
- [Boto3 S3 Documentation](#)
- [Python CSV Module](#)
- [Pika Documentation](#)

Deploying RabbitMQ on EC2

▶ Task:

- ▶ Launch a t2.micro EC2 instance to host RabbitMQ.

▶ Key Steps:

▶ Launch Instance:

- Choose Debian
- Select t2.micro
- Create a key pair, and assign an IAM instance profile (LabInstanceProfile).

▶ Configure Security Group:

- Allow traffic on ports 5672 (AMQP).

▶ Connect via SSH:

- Example command:

- ▶ `ssh -i "your-key.pem" ec2-user@<EC2_PUBLIC_IP>`



Deploying RabbitMQ on EC2

- ▶ **Install RabbitMQ and Dependencies**
 - ▶ Run the commands from the `ec2_steps.txt` file **on the EC2 instance you created.**
- ▶ **Resources:**
 - ▶ [Get started with AWS EC2](#)
 - ▶ [RabbitMQ Installation Guide](#)

Testing the Order Publisher Lambda

- ▶ **Action:**

- ▶ With RabbitMQ running, update your Lambda to include the RabbitMQ connection code.

- ▶ **Verification:**

- ▶ Re-upload the CSV file and check the RabbitMQ management console to see messages in the queue.



Delivery Simulator Lambda Function

► Objective:

- Create a Lambda function (named **lambda_delivery**) that simulates order delivery.

► Key Code Elements:

- **Extract Order:** Retrieve order details from the event.
- **Simulate Delay:** Use `random.randint(2, 10)` and `time.sleep(delay)` to simulate a delay.
- **Return Success:** Log the delivery and return a confirmation message.

► Important Note on Lambda Timeout:

- Make sure to configure your Lambda function's timeout to a sufficient value (for example, 15 seconds) in the AWS Lambda configuration.

► Resources:

- [Python `time.sleep\(\)`](#)
- [Random Module Documentation](#)

Building the Subscriber (Manager) on EC2

► Objective:

- Develop a Python client on an EC2 instance that subscribes to the RabbitMQ queue.
- For each received order, the Manager invokes the Delivery Simulator Lambda asynchronously.

► Key Code Components:

- **Connect to RabbitMQ:** Use pika to subscribe to the queue.
- **Lambda Invocation:** Use boto3 to call **lambda_delivery** with `InvocationType='Event'` (specify region as us-east-1).
- **Message Acknowledgement:** Ensure each message is acknowledged after processing.

► Resources:

- [Boto3 Lambda Client](#)
- [Pika Consumption Example](#)

Lab Summary and Final Verification

► Process Recap:

- Upload CSV file with orders to S3.
- Lambda (Order Publisher) processes the CSV and enqueues orders in RabbitMQ.
- A Manager (subscriber) running on EC2 reads orders from RabbitMQ.
- The Manager invokes the Delivery Simulator Lambda for each order.
- Each Delivery Lambda simulates order delivery in parallel.

► Verification:

- Monitor Lambda logs and RabbitMQ management console.
- Confirm that each order is processed and delivered concurrently.

Explore AWS SQS as an Alternative

- ▶ **Why AWS SQS?**
 - ▶ Fully managed, serverless messaging service.
 - ▶ No need to manage EC2 instances or clusters.
 - ▶ Integrates seamlessly with AWS Lambda.
- ▶ **Advantages Over RabbitMQ:**
 - ▶ Simplified operations & lower maintenance.
 - ▶ Automatically scales with demand.
 - ▶ Cost-effective—pay per use.
- ▶ **Your Challenge:**
 - ▶ Research AWS SQS documentation.
 - ▶ Consider how you'd adapt your messaging code to use SQS instead of RabbitMQ.
 - ▶ Explore SQS to enhance your lab experience and reduce operational overhead!

Additional Resources and Useful Commands

▶ AWS Documentation:

- [AWS Lambda Developer Guide](#)
- [AWS S3 User Guide](#)
- [AWS EC2 User Guide](#)



UNIVERSITAT
ROVIRA I VIRGILI

